



Boosting REST API Performance Using Asyncio

Johannes Krohn

Technical Marketing Engineer

10. Sep 2019

Agenda

Code, examples available at <https://github.com/jeokrohn/duasync>

- Python Performance
- Concurrency, ...
- `concurrent.futures`
- Asyncio
- Summary



Python Performance

Python Performance

- Depending on the workload Python code can be slower than other language like for example Java
- A lot of workloads are not really CPU bound
 - Other factors: network, queries, ...
 - Faster running code does not really help
- Python execution speed does not really matter when working on I/O bound problems

*Demo – Scraping “Spiegel”
sequentially*

I/O is very slow

- Each REST call takes "forever"
 - Creating the space
 - Adding a user to the space
- Python code the majority of the time just idles and needs to wait for a response

Concurrency, ...

Practical Example: Chess Exhibition (serial)

- Assumption:
 - 24 opponents
 - Master moves in 5 seconds
 - Other players move in 55 seconds
 - Game averages at 30 move pairs
- Each game runs 30 minutes
- 24 sequential games: 12 hours

Asynchronous Python for the Complete Beginner, Miguel Grinberg, Pycon 2017: <https://www.youtube.com/watch?v=iG6fr81xHKA>

Practical Example: Chess Exhibition (async)

- Master moves on first game
- Then moves to 2nd, 3rd, 4th, ..
- A move on all 24 games takes the master: $24 \times 5 \text{ sec} = 2 \text{ min}$
- After two minutes the 1st game is again ready for her move
- 24 games are completed in $2 \text{ min} \times 30 = 1 \text{ h}$

Asynchronous Python for the Complete Beginner, Miguel Grinberg, Pycon 2017: <https://www.youtube.com/watch?v=iG6fr81xHKA>


Parallelism, Concurrency, ...

- Parallelism: performing multiple operations at (virtually) the same time
 - Multiprocessing: spreading tasks over CPUs/cores
 - Good for CPU bound tasks
- Concurrency: multiple tasks can run in an overlapping manner
 - Does not imply parallelism
- Threading: concurrent execution model, threads take turns
 - Better for I/O bound tasks
- Python package: `concurrent.futures`

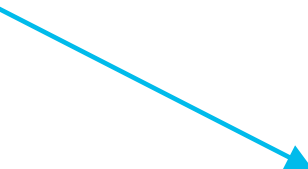
Demo – Scraping “Spiegel”
concurrent.futures

Findings: concurrent.futures

- ProcessPoolExecutor:
 - One Python process per task
 - Single thread per process
 - More memory intensive
- ThreadPoolExecutor:
 - Single Python process for all tasks
 - One thread per task



Process Name	Memory	Threads
python3.7	4,2 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	4,3 MB	1
python3.7	54,4 MB	4
python3.7	4,3 MB	1



Process Name	Memory	Threads
python3.7	53,4 MB	12

Asyncio Programming

- Single-threaded, single-process
- Cooperative multitasking
- “feeling of concurrency”
- What does “asynchronous” mean?
 - Async code can “pause” (wait for a result or events) and let other code run
 - Quasi concurrent execution of multiple tasks
 - “Cooperative”: execution is passed on when one task pauses

Asyncio Programming in Python

- Suspend and Resume
- Event loop: keep track of all async functions & state; schedule execution
- Options
 - Callbacks
 - Generators or coroutine functions
 - **Async/await**
- Options: **Asyncio**, Twisted, Tornado, ...

Blocking Library Functions

- Blocking calls are incompatible w/ async programming
 - `socket.*`, `select.*`
 - `threading.*`
- Async frameworks like (asyncio) need to provide replacements for these
- Fallback: run sync code in separate thread or process pool (`concurrent.futures`)
- No async support for file I/O

Asyncio

Asyncio: One Way for Async Python Programming

- Coroutines: `async` functions
- `await`: wait for result (pause) → execution can be passed to other coroutine
- Event loop: keep track of coroutines, state, scheduling
- Schedule tasks:
 - `loop.create_task()`
 - `asyncio.ensure_future()`
 - `asyncio.create_task()` – Python 3.7+ only

Demos: async1, async2, async3

aiohttp: Async Web Clients

- asyncio web client
- Replacement for synchronous `requests` library
- High-Level: `aiohttp.ClientSession()` object
 - Connection pooling
 - Cookie jar
 - Keepalive
 - ...
- Low-Level: `aiohttp.request()`

*Demo: Scraping “Spiegel” w/
asyncio*

Finding: `asyncio/aiohttp`

- Creating all tasks at the same time:
 - HTTP GET requests sent for all URLs
 - Issues:
 - Servicing the responses takes too long → server timeouts
 - Too many open files
 - Need to throttle the requests (similar to `max_workers` with `concurrent.futures`)
 - For example: semaphore

Applying `asyncio` to REST APIs

- Using `aiohttp` enables asynchronous REST calls
- Simply need to create the required calls using `aiohttp` methods

*Demo: Webex Teams demos
(sync vs async)*

Limitation of “Naive” REST API Approach

- No readily available replacement for sync libraries
 - Each endpoint has to be coded explicitly
- Webex Teams specific:
 - Support for 429 rate limiting
 - Pagination
 - ...
- Potential solution: use “non asyncio” library and `concurrent.futures` with `ThreadPoolExecutor`

Summary

What Form of Concurrency is “best”?

```
if io_bound:
    if io_very_slow or lots_of_connections:
        print('use asyncio')
    else:
        print('use threads')
else:
    print('use processes')
```

Problems with asyncio

- Need to think “asyncio”
- Hard to migrate existing code
- Some sync libraries don't have async equivalence (for example `webexteamssdk`)

Other Services

- FTP (server/client): [aioftp](#)
- Timeouts: [async-timeout](#)
- SSH: [asyncssh](#)
- Client WebSockets: [aiohttp](#)
- Interaction with network devices: [netdev](#)
- Others: <https://github.com/aio-libs>

References

- <https://realpython.com/async-io-python/>
- <https://realpython.com/python-concurrency/>
- https://training.talkpython.fm/courses/explore_async_python/async-in-python-with-threading-and-multiprocessing
- <https://docs.python.org/3/library/asyncio.html>
- <https://docs.python.org/3/library/concurrent.futures.html>

